

# Integrating Ecovisor into Mosaik Co-Simulation

Marvin Steinke and Henrik Nickel  
*Technische Universität Berlin*

**Abstract**—To reduce emissions, cloud platforms must increasingly rely on renewable energy sources such as solar and wind power. Nevertheless, the issue of volatility associated with these sources presents a significant challenge, since current energy systems conceal such unreliability in hardware. Souza et al. have devised a solution to this issue by creating an “ecovisor”. This system virtualizes the energy infrastructure and allows for software-defined control to be accessible by applications.

## I. INTRODUCTION

THE surge of cloud platforms has had a significant impact on many businesses and individuals, offering access to innovative and valuable applications that frequently require significant computational resources. A notable example for this is represented by OpenAI’s chat generative pre-trained transformer (ChatGPT) cloud platform, which is based on the GPT-3 model developed by Brown et al. [1] and achieved more than a million users within the first five days of its launch [2]. The increasing demand for more computational power has lead cloud platforms to become an essential part of the digital landscape [3]. These platforms allow for the storage, processing and management of large amounts of data and computational resources, which can be used to run applications and services that are beyond the capabilities of traditional hardware systems.

However, while the growth of cloud platforms has brought many benefits, it has also raised concerns about their environmental impact. As the demand for computational resources increases, so does the carbon emissions generated by the energy consumption of these platforms [4]. Despite their ecological impact, the growth of cloud platforms shows no signs of slowing down. According to Gartner Inc. worldwide end-user spending on public cloud services is forecasted to grow 20.7% to total \$591.8 billion in 2023 [5]. To mitigate their impact on the environment, cloud platforms are now looking for ways to reduce their carbon footprint [6, 7]. It is imperative to adopt cleaner energy sources for powering data centers, both in the cloud and at the edge.

Although clean energy offers numerous advantages, it is perceived as being unreliable due to two key factors. One, the generation of renewable energy sources such as solar and wind is affected by environmental changes, and two, the carbon-intensity of grid power fluctuates as the grid uses different types of generators with varying carbon emissions to meet demand [8]. The field of computing possesses a distinct advantage in terms of reducing its carbon impact through the utilization of cleaner energy sources [9]. However, current cloud applications are unable to apply these benefits to optimize their carbon efficiency. This is because the energy system obscures the instability of clean energy with a

reliability abstraction, which gives no control or visibility into the energy supply. As a result, applications cannot adjust their power usage in response to changes in the availability and carbon-intensity of renewable energy [10].

Souza et al. [10] proposed a solution to this issue by creating an “ecovisor”, which virtualises the energy system and provides software-based control over it. This ecovisor enables each application to manage the instability of clean energy through software, customized to meet its unique requirements. However, even their small-scale prototype is intricately designed and incorporates several expensive components, such as a DC power supply equipped with a solar array simulation capability that costs almost \$10,000.

Though the creation of the ecovisor is a promising solution to the issue of reducing the carbon footprint of cloud platforms, a large scale prototype for research and development purposes would not only be significantly cost intensive, but also time consuming. An alternative approach to consider is to simulate *only* the ecovisor infrastructure with real applications as a Software-In-The-Loop (SIL) implementation. This would allow for the optimization of energy usage in response to changes in the availability and carbon-intensity of renewable energy, without the need for a physical implementation. By simulating the ecovisor, applications can still manage the instability of clean energy through software, but at a lower cost and with less time investment.

To this end, we propose the utilization of a co-simulation tool, Mosaik, to integrate this ecovisor simulation into, and evaluate its impact on the carbon footprint of cloud platforms. Mosaik is a simulation framework for power systems, communication networks, and building automation, which can be used to model and analyze the performance of the ecovisor in real-world applications [11].

In the following sections of this article, we will commence by presenting the necessary contextual information in Section II to ensure comprehension of our methodology. This includes co-simulation environments with Software- (SIL) and Hardware-In-The-Loop (HIL) strategies, the Mosaik co-simulation tool and a closer examination of the ecovisor infrastructure and its interface to applications. We will then review some related literature with similar approaches in Section III. Subsequently, we present our approach of integrating the ecovisor into Mosaik in Section IV. The approach is divided into the simulation of the ecovisor itself and the interface to external applications, beyond the scope of the co-simulation. We evaluate our approach in Section V with exemplary data, representing a realistic scenario that covers all edge cases. After discussing future research directions related to this article in Section VI, we conclude this article in Section VII

by summarizing the main points and contributions.

## II. BACKGROUND

### A. Co-Simulations and SIL / HIL

Co-simulations are used to model and analyze complex systems with multiple interacting components, each of which may have different properties and behaviors. They involve combining simulation models from different domains, such as control, power, and thermal management, to create a unified model that accurately represents the behavior of the overall system.

One of the main advantages of co-simulations over regular simulations is their ability to capture the interactions between different components of the system. Regular simulations often make simplifying assumptions that can lead to inaccurate results. Co-simulations, on the other hand, can account for the interactions between components and provide a more accurate representation of the system's behavior. This makes co-simulations particularly useful for designing and optimizing complex systems like datacenters. The virtual environment can save time and resources, reduce the risk of failure, and lead to more efficient and sustainable datacenters [12].

Co-simulations can be further enhanced by incorporating the SIL and HIL methodology. This approach involves integrating real-world components, such as software and hardware, into the simulation environment to better reflect the actual system behavior. The integration of real components allows for a more accurate representation of the system's behavior and can also identify potential issues that may arise in real-world scenarios [13].

### B. Mosaik

Mosaik is an open-source co-simulation tool that allows for the integration of different simulation models from various domains into a unified simulation environment [11]. Mosaik's four main components enable communication between simulators and Mosaik, the creation of simulation scenarios, the management of simulator processes, and the coordinated simulation of a scenario.<sup>1</sup>

**Mosaik Sim API.** The Mosaik Sim API provides a standardized communication protocol for simulators and Mosaik to exchange information. It utilizes plain network sockets and JSON encoded messages to facilitate communication between simulators and Mosaik. Although the API provides a low-level communication protocol, it is complemented by a high-level API for some programming languages. The high-level API abstracts the networking-related aspects of the communication protocol, allowing users to focus on the core logic of their simulation models. To use the high-level API, users only need to write a subclass and implement a few required methods.

**Scenario API.** This empowers users to create simulation scenarios entirely in Python. This API enables users to launch simulators, create model instances, and generate

entity sets that can be connected to establish data flows between different simulators. With the Scenario API, users can easily connect individual entities or entire sets of entities to achieve their simulation goals.

**SimManager.** The Simulator Manager, or SimManager, handles simulator processes and their communication with Mosaik. The SimManager can start new simulator processes, connect to already running process instances, and import a simulator module and execute it in-process if it is written in Python 3. In-process execution reduces memory requirements and avoids the overhead of serializing and sending messages over the network. External processes, however, can be executed in parallel, which is not possible with in-process simulators.

**Mosaik's Scheduler.** Mosaik's Scheduler uses the event-discrete simulation library SimPy for the coordinated simulation of a scenario. Mosaik supports time-discrete and event-discrete simulations, as well as a combination of both paradigms. It can handle simulators with different step sizes, and a simulator may vary its step size during the simulation. Mosaik tracks the dependencies between simulators and only lets them perform a simulation step if necessary. It is also able to let multiple simulators perform their simulation step in parallel if they do not depend on each other's data.

### C. Ecovisor<sup>2</sup>

Figure 1 shows an overview of the ecovisor's design which manages resources and energy using containers or virtual machines as the basic unit. An instance-level API is chosen to align with existing cloud APIs and to support higher-level cluster or cloud-level APIs. The ecovisor extends an existing orchestration platform that provides basic container or VM management and monitoring functions. Container Orchestration Platforms (COPs) are used to manage resources and applications. COPs provide virtual clusters composed of multiple containers with specified resource allocations that can grow or shrink over time. COPs include a scheduling policy that determines resource allocation under constraints, and COPs are resilient to resource revocations. This resiliency is useful for designing carbon-efficient applications as low-carbon energy may cause power shortages that also manifest as resource revocations.

The virtual energy system includes a virtual grid connection, a virtual battery, and a virtual solar array. The system provides getters and setters methods for monitoring and controlling the virtual power supply and demand, including per-container power caps, battery charging, and discharging rates, as shown in Table I. The system uses virtual solar power first to meet demand and charges the virtual battery with any excess solar power. When there is not enough solar power, the virtual energy system uses grid power to make up the difference, while accounting for carbon emissions and power usage over discrete time intervals. The ecovisor system provides a uniform

<sup>1</sup>The information about the main components is adapted from the official documentation [14].

<sup>2</sup>The information presented in this section is a summary of Section 3 and 3.1 from Souza et al.'s work [10], with some modifications made for clarity and conciseness.

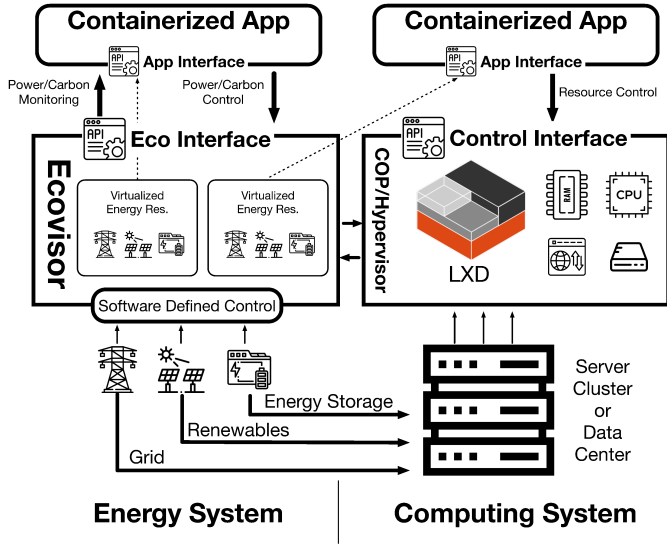


Figure 1. Ecovisor Design (Souza et al.) [10]

centralized interface for accessing energy-related information and historical data.

### III. RELATED WORK

In research and development processes, the utilization of physical testbeds is a widespread practice [15, 16]. They serve as a medium for testing and evaluating new technologies, systems, and products prior to their market launch or further development. However, some scenarios and conditions may not be feasible or safe to recreate in a physical testbed. Simulations can provide a controlled and cost-effective environment for testing such scenarios and conditions [17]. When specific physical or software components require a certain degree of realism though, simulations or co-simulations with SIL or HIL methodologies can effectively address these limitations.

For instance, Beilharz et al. [18], introduce Marvis, a framework that provides a comprehensive staging environment for testing distributed IoT applications. Marvis orchestrates hybrid testbeds, co-simulated domain environments, and a central network simulation to create a representative operating environment for the system. However, Marvis does not provide a virtualized energy system, which is crucial to meet the requirements of our problem statement.

Hagenmeyer et al. [19] investigate the interplay of different forms of energy on various value chains in Energy Lab 2.0. The focus is on finding novel concepts to stabilize the volatile energy supply of renewables through the use of storage systems and the application of information and communication technology tools and algorithms. The smart energy system simulation and control center is a key element of Energy Lab 2.0 and consists of three parts: a power-hardware-in-the-loop experimental field, an energy grid simulation and analysis laboratory, and a control, monitoring, and visualization center. While this smart energy system simulation is similar to the simulated ecovisor in our approach, the control center is the only entity with software-based control over this energy system. The ecovisor infrastructure, however, provides multiple

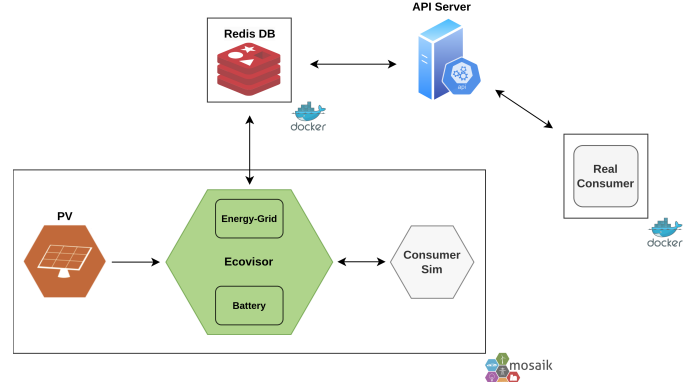


Figure 2. General System Design: The ecovisor infrastructure is simulated and integrated into the Mosaik co-simulation environment. SIL capabilities are enabled via an API Server and a Redis database, connecting the simulation environment and real applications in real-time.

applications with the ability to manage their energy supply themselves which is an essential factor for our approach.

In conclusion, to the best of our knowledge, there is no current approach that can simulate an energy system virtualizer with software-based control, comparable to the ecovisor, with SIL capabilities.

### IV. APPROACH

Figure 2 illustrates our general system design approach that simulates the ecovisor infrastructure and integrates it into the Mosaik co-simulation environment while enabling SIL capabilities. The present design can be categorized into two distinct parts, the simulation of the ecovisor, and its interface to external applications, which are elaborated on in the following.

#### A. Simulation of the Ecovisor

The simulation part of our approach is fully contained within the Mosaik co-simulation environment. In Figure 2 this includes the photo-voltaic (PV) module, the ecovisor and a consumer simulator. The consumer and PV module simulators utilized in this context are the `mosaik-csv` simulator, which is a component of the Mosaik ecosystem [20]. This simulator is capable of simulating CSV datasets, and in the current application, it is employed to simulate consumption data and recorded or forecasted PV data. While our primary focus is to facilitate the integration of external, non-simulated applications with our simulated ecovisor, we also recognize the need to accommodate simulated consumers for more sophisticated approaches that require additional data. Therefore, we aim to provide flexibility in our system to meet a wide range of requirements. Because these two components are designed to be minimal and replaceable, they each require a controller agent that acts as a medium between the component and the simulated ecovisor. The controller can be initialized with e.g. a power conversion factor to allow for the use of different measurement units in the datasets, as the standard power unit utilized in the ecovisor is kW.

Table I  
ECOVISOR'S API (SOUZA ET AL.) [10]

Function Name	Type	Input	Return Value	Description
set_container_powercap()	Setter	ContainerID, kW	N/A	Set a container's power cap
set_battery_charge_rate()	Setter	kW	N/A	Set battery charge rate until full
set_battery_max_discharge()	Setter	kW	N/A	Set max battery discharge rate
get_solar_power()	Getter	N/A	kW	Get virtual solar power output
get_grid_power()	Getter	N/A	kW	Get virtual grid power usage
get_grid_carbon()	Getter	N/A	g · CO <sub>2</sub> /kW	Get virtual grid power usage
get_battery_discharge_rate()	Getter	N/A	kW	Get current rate of battery discharge
get_battery_charge_level()	Getter	N/A	kWh	Get energy stored in virtual battery
get_container_powercap()	Getter	ContainerID	kW	Get a container's power cap
get_container_power()	Getter	ContainerID	kW	Get a container's power usage
tick()	Notification	N/A	N/A	Invoked by ecovisor every $\Delta t$

```

1 rest ← consumption − solar
2 if rest ≤ 0 then
3   | b_discharge_rate ← 0
4 else
5   | b_discharge_rate ← min(
6     | b_max_discharge,
7     | b_charge_level · 3600,
8     | rest
9   | )
10  rest ← rest − b_discharge_rate
11 end
12 grid_power ← b_charge_rate + rest
13 b.delta ← b_charge_rate − b_discharge_rate
14 b.step()
15 b_charge_level ← b.charge
16 total_carbon ← grid_carbon · grid_power

```

Algorithm 1: Virtual Energy System Simulation

### B. Interface to External Applications

To connect the Ecovisor-Model to a real workload-application, we exposed the API which is described in [10] to containerized workloads. We have tried to implement it in a way that is as close as possible to a real implementation in a cloud environment. To achieve that, we implemented the Ecovisor-API from [10], which can be seen in I, into an FastAPI server [21]. The API-Server is connected to a RedisDB[22], which serves as a key-value store and links the API-Server to the Ecovisor-Model. The Ecovisor-Model itself implements an Redis-Interface to read and write data from the RedisDB. In the following three sections we will describe these three parts of the *External Application Interface* in depth.

1) *API-Server*: The API-Server exposes the Ecovisor-API to the "Workloads". It is implemented with the FastAPI Framework and utilizes the Uvicorn Web-Server [23] to handle multiple clients accessing the API. Due to the ASGI (Asynchronous Server Gateway Interface) nature of the Web-Server, the module is started as an independent thread. The advantage of this, is that the server can handle multiple clients, which will be useful when multiple workloads will be connected to the simulation. On the other hand, the implementation into the system is more difficult, so that we decided to start

```

1 {
2   "solar_power" : "0kW",
3   "grid_power" : "0kW",
4   "grid_carbon" : "0g * CO2/kW",
5   "battery_discharge_rate" : "0kW",
6   "battery_charge_level" : "0kWh"
7 }

```

Algorithm 2: Example Redis Data JSON

the module independently. In earlier versions we tried to implement it inside of the Ecovisor-Model, but it would stop the execution of the simulation until the API-Server is stopped which renders the system unusable. This also enables the API-Server to be scaled independent from the Ecovisor-Model and the RedisDB. This may be helpful in bigger simulations with distributed infrastructure.

2) *RedisDB*: The RedisDB is a fast, in memory, key-value store for simple data-types. In our system, it is used to save and provide the data from within the Ecovisor-Model to the API-Server. To keep the operability of the simulation simple, the RedisDB is started as a Docker Container via the docker python library [24], which implements the docker engine API into python. The container image we used is the redislabs/rejson redis image [25]. This image extends the RedisDB image with the capability to process JSON values. The Redis-Container is started right at the beginning of the simulation, so that it is up and operational when the first data is available. After the end of the simulation, the container is stopped and deleted to keep the test environment free of remnants of the simulation. In general the database can easily be exchanged with any other database by adapting the Ecovisor-Redis-Interface in the Ecovisor-Model. This can be useful when simulation should be integrated in a production grade cloud environment like kubernetes or openstack.

3) *Ecovisor-Redis-Interface*: The interface from the ecovisor towards the redisdb is implemented within the ecovisor model. To simplify the usage, the redis python library is used. The ecovisor model implements two methods, the `get_redis_updates()` and the `send_redis_updates`

The `get_redis_update` method gets the values (...) as




Figure 3. Dataflow between Ecovisor and API Server

show in (ref) from the RedisDB and writes it to the ecovisor model variables to make them usable for the ecovisor model.

the `send_redis_updates()` method publishes the updated values from the other part of the simulation which are accessible by the ecovisor shown in (ref) model to the redisdb to make it usable by the workload application.

In each step the ecovisor model first gets the updates from the RedisDB, processes the data and then sends the updated data to the RedisDB, to make it accessible to the workload application.

4) *Dataflow*: The Values/Data shown in (ref) is exchanged between the ecovisor model and the API server via the RedisDB. There are two main streams of the dataexchange. First from the ecovisor model to the API-Server and second from the API-Server to the ecovisor model. There are no values which are used in both streams, which makes things easier in terms of consistency and parallel access.

As stated before, the data is updated in each simulation step. (ref) In the downstream, from the ecovisor to the API-Server the values (...) are updated and can be consumed by the workload applications and in the upstream, from the api-server to the ecovisor model, the values (...) are updated and used by the ecovisor model to calculate (...)

## V. EVALUATION

## VI. FUTURE WORK

While our approach focused on a single simulation of the ecovisor, our use of Mosaik demonstrates its effectiveness for large-scale smart grid simulations. In future research, we suggest interconnecting multiple ecovisor systems to share resources and further reduce carbon emissions. This network could be distributed across different geographic regions, as carbon intensity varies depending on location. By incorporating carbon information services like Electricity Maps [26], this could enable carbon-efficiency optimizations such as Let's Wait Awhile [27] or Cucumber [28] from Wiesner et al. This would enhance the potential for carbon reduction at a larger scale.

## VII. CONCLUSION

## REFERENCES

[1] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell *et al.*, "Language mod-

els are few-shot learners," *Advances in neural information processing systems*, vol. 33, pp. 1877–1901, 2020.

[2] G. Brockman, "Chatgpt just crossed 1 million users," Twitter, dec 2022, accessed: 16.02.2023. [Online]. Available: <https://twitter.com/gdb/status/1599683104142430208?s=20&t=-7V8elwpZ93qilhwznzVBA>

[3] Q. Zhang, L. Cheng, and R. Boutaba, "Cloud computing: state-of-the-art and research challenges," *Journal of internet services and applications*, vol. 1, pp. 7–18, 2010.

[4] M. A. B. Siddik, A. Shehabi, and L. Marston, "The environmental footprint of data centers in the united states," *Environmental Research Letters*, vol. 16, no. 6, p. 064017, 2021.

[5] Gartner Inc., "Gartner forecasts worldwide public cloud end-user spending to reach nearly \$600 billion in 2023," oct 2022, accessed: 16.02.2023. [Online]. Available: <https://www.gartner.com/en/newsroom/press-releases/2022-10-31-gartner-forecasts-worldwide-public-cloud-end-user-spending-to-reach-nearly-600-billion-in-2023>

[6] Google LLC, "Building a carbon-free future for all," accessed: 16.02.2023. [Online]. Available: <https://sustainability.google/commitments/carbon/>

[7] Amazon.com, Inc., "Sustainability in the cloud," accessed: 16.02.2023. [Online]. Available: <https://sustainability.aboutamazon.com/environment/the-cloud>

[8] N. Scarlat, M. Prussi, and M. Padella, "Quantification of the carbon intensity of electricity produced and used in europe," *Applied Energy*, vol. 305, p. 117901, 2022.

[9] S. Murgesan, "Harnessing green it: Principles and practices," *IT Professional*, vol. 10, no. 1, pp. 24–33, 2008.

[10] A. Souza, N. Bashir, J. Murillo, W. Hanafy, Q. Liang, D. Irwin, and P. Shenoy, "Ecovisor: A virtual energy system for carbon-efficient applications," in *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*, 2023, pp. 252–265.

[11] C. Steinbrink, M. Blank-Babazadeh, A. El-Ama, S. Holly, B. Lüers, M. Nebel-Wenner, R. P. Ramírez Acosta, T. Raub, J. S. Schwarz, S. Stark, A. Nieße, and S. Lehnhoff, "Cpes testing with mosaik: Co-simulation planning, execution and analysis," *Applied Sciences*, vol. 9, no. 5, 2019.

[12] M. Vogt, F. Marten, and M. Braun, "A survey and statistical analysis of smart grid co-simulations," *Applied Energy*, vol. 222, pp. 67–78, 2018.

[13] T. Kelemenová, M. Kelemen, L. Miková, V. Maxim, E. Prada, T. Lipták, and F. Menda, "Model based design and hil simulations," *American Journal of Mechanical Engineering*, vol. 1, no. 7, pp. 276–281, 2013.

[14] OFFIS – Institut für Informatik, "Mosaik 3.1.1 documentation," accessed: 16.02.2023. [Online]. Available: <https://mosaik.readthedocs.io>

[15] M. H. Cintuglu, O. A. Mohammed, K. Akkaya, and A. S. Uluogac, "A survey on smart grid cyber-physical system testbeds," *IEEE Communications Surveys & Tutorials*, vol. 19, no. 1, pp. 446–464, 2017.

[16] J. Mambretti, J. Chen, and F. Yeh, "Next generation clouds, the chameleon cloud testbed, and software defined networking (sdn)," in *2015 International Conference on Cloud Computing Research and Innovation (ICCCRI)*, 2015, pp. 73–79.

[17] N. Mansouri, R. Ghafari, and B. M. H. Zade, "Cloud computing simulators: A comprehensive review," *Simulation Modelling Practice and Theory*, vol. 104, p. 102144, 2020.

[18] J. Beilharz, P. Wiesner, A. Boockmeyer, F. Brokhausen, I. Behnke, R. Schmid, L. Pirl, and L. Thamsen, "Towards a staging environment for the internet of things," in *2021 IEEE International Conference on Pervasive Computing and Communications Workshops and other Affiliated Events (PerCom Workshops)*, 2021, pp. 312–315.

[19] V. Hagenmeyer, H. Kemal Çakmak, C. Döpmeier, T. Faulwasser, J. Isele, H. B. Keller, P. Kohlhepp, U. Kühnapfel, U. Stucky, S. Waczowicz, and R. Mikut, "Information and communication technology in energy lab 2.0: Smart energies system simulation and control center with an open-street-map-based power flow simulation example," *Energy Technology*, vol. 4, no. 1, pp. 145–162, 2016.

[20] OFFIS – Institut für Informatik, "mosaik-components contains some components, which can be used as base for a co-simulation scenario," accessed: 16.02.2023. [Online]. Available: <https://gitlab.com/mosaik/components>

[21] Sebastián Ramírez, "Fastapi," accessed: 03.03.2023. [Online]. Available: <https://fastapi.tiangolo.com>

[22] Redis Ltd, "Redis," accessed: 03.03.2023. [Online]. Available: <https://redis.io>

[23] Encode UK Ltd, "Uvicorn," accessed: 03.03.2023. [Online]. Available: <https://www.uvicorn.org>

[24] Docker Inc, "Python docker," accessed: 03.03.2023. [Online]. Available: <https://github.com/docker/docker-py>

- [25] Redis Ltd, “Redisjson,” accessed: 03.03.2023. [Online]. Available: <https://hub.docker.com/r/redislabs/rejson>
- [26] Electricity Maps A/S , “Reduce carbon emissions with actionable electricity data,” accessed: 16.02.2023. [Online]. Available: <https://www.electricitymaps.com/>
- [27] P. Wiesner, I. Behnke, D. Scheinert, K. Gontarska, and L. Thamsen, “Let’s wait awhile,” in *Proceedings of the 22nd International Middleware Conference*. ACM, dec 2021.
- [28] P. Wiesner, D. Scheinert, T. Wittkopp, L. Thamsen, and O. Kao, “Cucumber: Renewable-aware admission control for delay-tolerant cloud and edge workloads,” in *Euro-Par 2022: Parallel Processing*. Springer International Publishing, 2022, pp. 218–232.